

Chapter 4

Application: Identifying Trending Topics on Twitter

In this chapter, we consider the application of the method and algorithm proposed in chapters 2 and 3 toward detection of trending topics on Twitter. we discuss the Twitter service, the collection and pre-processing of data, and the experimental setup for the detection task.

4.1 Overview

4.1.1 Overview of Twitter

Twitter is a real-time messaging service and information network. Users of Twitter can post short (up to 140 characters) messages called *Tweets*, which are then broadcast to the users' *followers*. Users can also engage in conversation with one another. By default, Tweets are public, which means that anyone can see them and potentially join a conversation on a variety of topics being discussed. Inevitably, some topics gain relatively sudden popularity on Twitter. For example, a popular topic might reflect an external event such as a breaking news story or an internally generated inside joke or game. Twitter surfaces such topics in the service as a list of Trending Topics.

4.1.2 Twitter-Related Definitions

Talking about Tweets, topics, trends and trending topics can be ambiguous, so here we make precise our usage of these and related terms.

Definition 1 (Topic). *We define a **topic** to be a phrase consisting of one or more words delimited by spacing or punctuation. A word may be any sequence of characters and need not be an actual dictionary word.*

Definition 2 (Tweet about topic). *A Tweet is **about** a topic if it contains the topic as a substring.*

Definition 3 (Trending topic). *A **trending topic** is a topic that is currently on the list of Trending Topics on Twitter. If a topic was ever a trending topic during a period of time, we say that the topic **trended** during that time period.*

Definition 4 (Trend). *A trending topic will also occasionally be referred to as a **trend** for short.*

Definition 5 (Trend onset). *The **trend onset** is the time that a topic first trended during a period of time.*

4.1.3 Problem Statement

At any given time there are many topics being talked about on Twitter. Of these, some will trend at some point in the future and others will not. We wish to predict which topics will trend. The earlier we can predict that a topic will trend, the better. Ideally, we would like to do this while maintaining a low rate of error (false detections and false non-detections).

4.1.4 Proposed Solution

Our approach to detecting trending topics is as follows. First, we gather examples of topics that trended and topics that did not trend during some period of time. Then, for each topic, we collect Tweets about that topic and generate a time series of the

activity of that topic over time. We then use those time series as *reference signals* (cf. Chapter 2) and apply the classification method and algorithm described in Chapters 2 and 3.

4.2 Data

4.2.1 Data Collection

The online time series classification method detailed in Chapters 2 and 3 requires a set of reference signals corresponding to trending topics and a set of reference signals corresponding to non-trending topics. These reference signals represent historical data against which we can compare our most recent observations to do classification.

The data collection process can be summarized as follows. First, we collected 500 examples of topics that trended at least once between June 1, 2012 and June 30, 2012 (hereafter referred to as the *sample window*) and 500 examples of topics that never trended during the sample window. We then sampled Tweets from the sample window and labeled each Tweet according to the topics mentioned therein. Finally, we constructed a reference signal for each topic based on the Tweet activity corresponding to that topic.

We obtained all data directly from Twitter via the MIT VI-A thesis program. However, the type as well as the amount of data we have used is all publicly available via the Twitter API.

Topics

We collected a list of all trending topics on Twitter from June 1, 2012 to June 30, 2012 (the *sample window*) as well as the times that they were trending and their rank in the Trending Topics list on Twitter. Of those, we filtered out topics whose rank was never better than or equal to 3. In addition, we filtered out topics that did not trend for long enough (the time of the first appearance to the time of the last appearance is less than 30 minutes) as well as topics that reappear multiple times during the

sample window (the time of the first appearance to the time of the last appearance is greater than 24 hours). The former eliminates many topics that are spurious and only trend for a very short time. The latter eliminates topics that correspond to multiple events. For example, the name of a football player might trend every time there is an important game. We would like to avoid such ambiguity and restrict each trending topic to correspond to a single underlying event within the sample window.

We collected topics that did not trend during the sample window in two steps. First, we sampled a list of n -grams (phrases consisting of n “words”) occurring on Twitter during the sample window for n up to 5. We filtered out n -grams that contain any topic that trended during the sample window, using the original, unfiltered list of all topics that trended during the sample window. For example, if “Whitney Houston” trended during the sample window, then “Whitney” would be filtered out of the list of topics that did not trend. We also removed n -grams shorter than three characters, as most of these did not appear to be meaningful topics. Lastly, we sampled 500 n -grams uniformly from the filtered list of n -grams to produce the final list.

Tweets

We sampled 10% of all public Tweets from June 1, 2012 to June 30, 2012 inclusive. We labeled each Tweet with the topic or topics contained therein using a simple regular expression match between the Tweet text and the topic text. In addition to the Tweet text, we recorded the date and time the Tweet was authored.

4.2.2 From Tweets to Signals

We discuss the process of converting the timestamped Tweets for a given topic into a reference signal. Each of the steps below is followed in order for each topic.

Computing Tweet Rate

First, we determine the rate of Tweets about a topic over time by binning the Tweets into time bins of a certain length. we used time bins of length two minutes.

Normalizing to Remove Baseline

A first glance at the data reveals that many non-trending topics have a relatively high rate and volume of Tweets, and many trending topics have a relatively low rate and volume of Tweets. One important difference is that many non-trending topics have a high *baseline rate* of activity while most trending topics are preceded by little, if any, activity prior to gaining sudden popularity. For example, a non-trending topic such as 'city' is likely to have a consistent baseline of activity because it is a common word. To emphasize the parts of the rate signal above the baseline and de-emphasize the parts below the baseline, we divide the rate by the mean rate over the entire window. We then take the resulting quotient to a power β , which controls how much we reward and penalize rates above and below the mean rate. Figure ?? shows signals without any baseline normalization and their baseline-normalized versions.

Rewarding Spikes

Another difference between the rates of Tweets for trending topics and that of non-trending topics is the number and magnitude of spikes. The Tweet rates for trending topics typically contains larger and more sudden spikes than that of non-trending topics. We reward such spikes by emphasizing them, while de-emphasizing smaller spikes. If the signal so-far in the transformation process is r , we transform it to

$$s(t) = |\dot{r}(t)|^\alpha,$$

or in the discrete case,

$$s[n] = |r[n] - r[n-1]|^\alpha,$$

where $\alpha > 1$. Figure ?? shows the effect of this spike-based transformation.

Slicing

The signal resulting from the steps so far is as long as the entire time window from which all Tweets were sampled. Such a long signal is not particularly useful as a reference signal. Recall from chapter 3 that to see how much the recent trajectory of the observed signal resembles part of a reference signal, we have to traverse the full length of the reference signal in order to find the piece that most closely resembles the recent observed trajectory. If the reference signal for topic that trended spans too long of a time window, only a small portion of it will represent activity surrounding the onset of the trend. In addition, it is inefficient to compare the recently observed trajectory to a reference signal that is needlessly long. Hence, it is necessary to select a small slice of signal from the much longer rate signal. In the case of topics that trended, we select a slice that terminates at the first onset of trend. That way, we capture the pattern of activity leading up to the trend onset, which is crucial for recognizing similar pre-onset activity in the observed signal. We do not include activity after the true onset because once the topic is listed in the Trending Topics on Twitter, we expect the predominant mode of spreading to change. For topics that did not trend, we assume that the rate signal is largely stationary and select slices with random start and end times. For simplicity, all slices are a fixed size.

Smoothing

Tweet rates, and the aforementioned transformations thereof, tend to be noisy, especially for small time bins. To mitigate this, we convolve the signal resulting from the previous step with a smoothing window of size N_{smooth} . Figure 4-2 shows the effect of smoothing with various window sizes.

Branching Processes and Logarithmic Scale

It is reasonable to think of the spread of a topic from person to person as a branching process. A branching process is a model of the growth of a population over time, in which each individual of a population in a given generation produces a random

number of individuals in the next generation. While we do not know the details of how a topic spreads, we do know that in a wide generality of branching processes, the growth of the population is exponential with time, with the exponent depending on the details of the model [?] (see [?] for a treatment involving the spread of topics on Twitter). It is reasonable, then, to measure the volume of tweets at a logarithmic scale to reveal these details. Therefore, as a final step, we take the logarithm of the signal constructed so far. Figure 4-3 shows a sample of signals and their log-scaled versions.

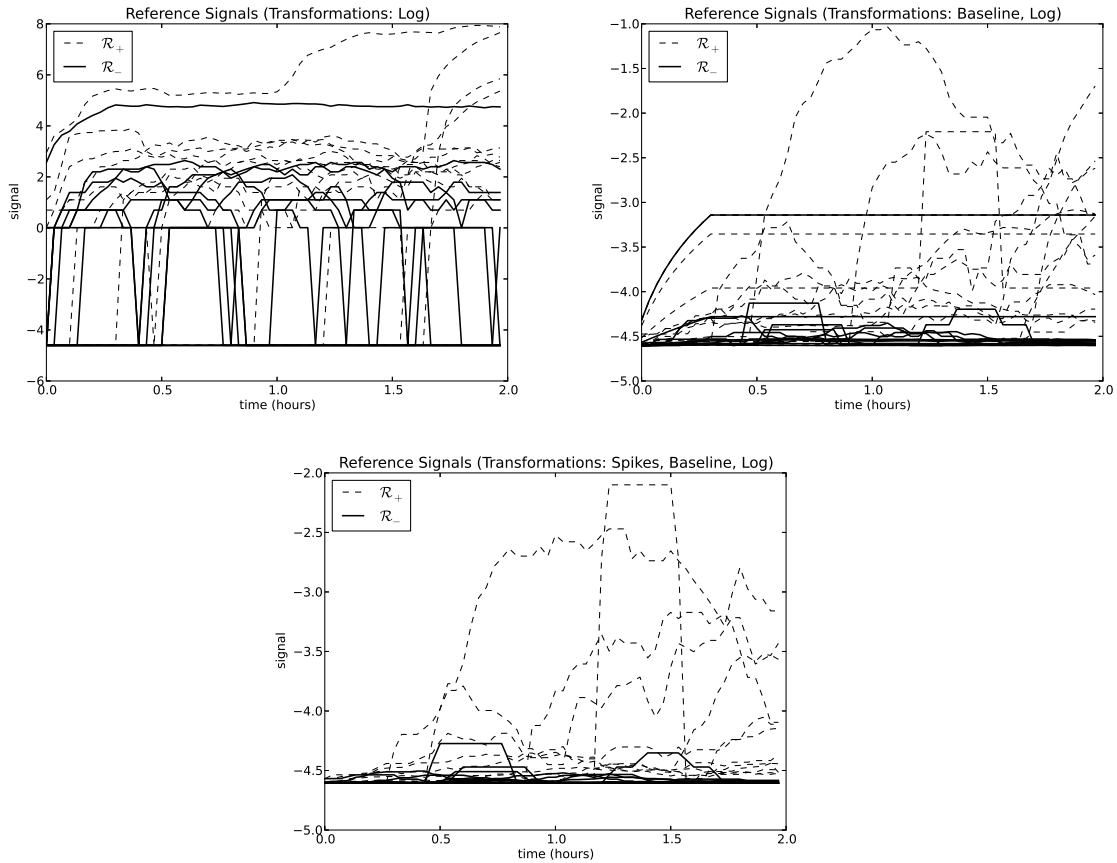


Figure 4-1: Reference signals of either class are hard to tell apart without normalization. **Top left:** no baseline or spike normalization. **Top right:** Baseline normalization. **Bottom:** Baseline and spike-based normalization.

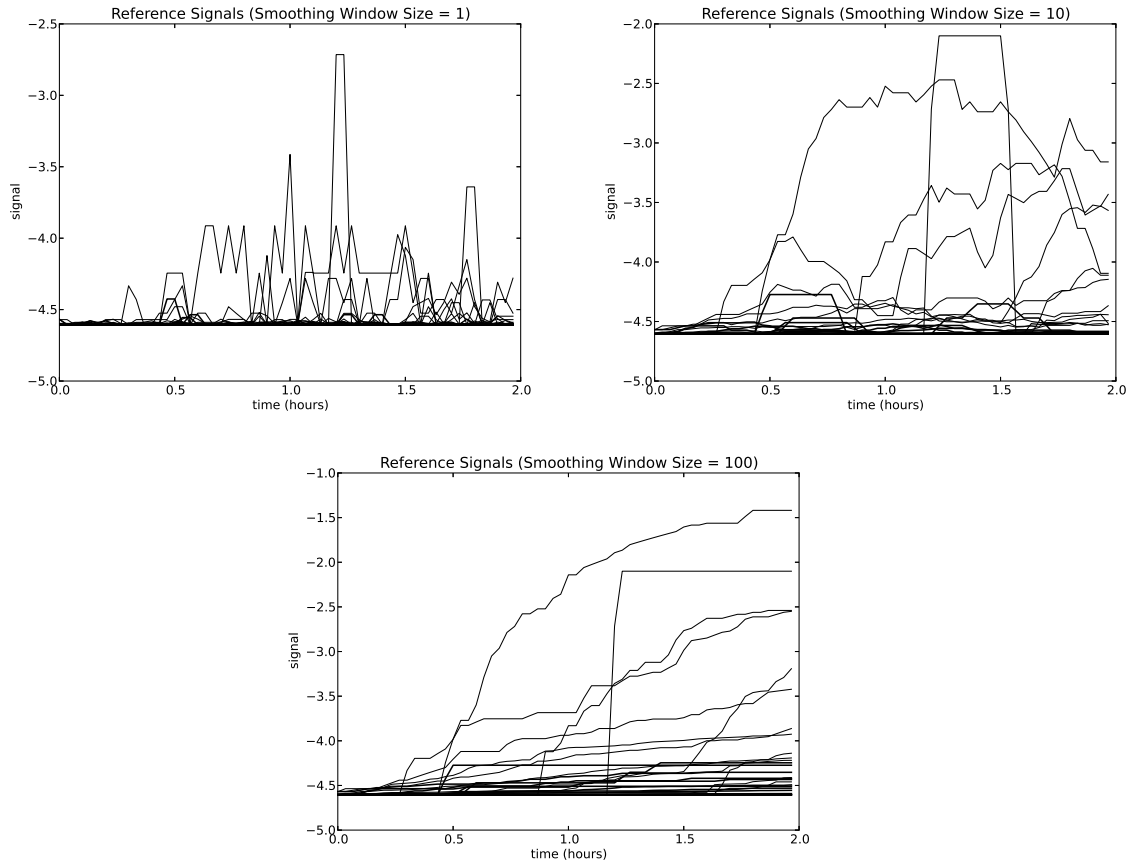


Figure 4-2: The results of smoothing the reference signals (with spike and baseline normalization previously applied) with windows of size 1 (2 minutes, i.e. no smoothing), 10 (20 minutes), and 100 (3 hours, 20 minutes).

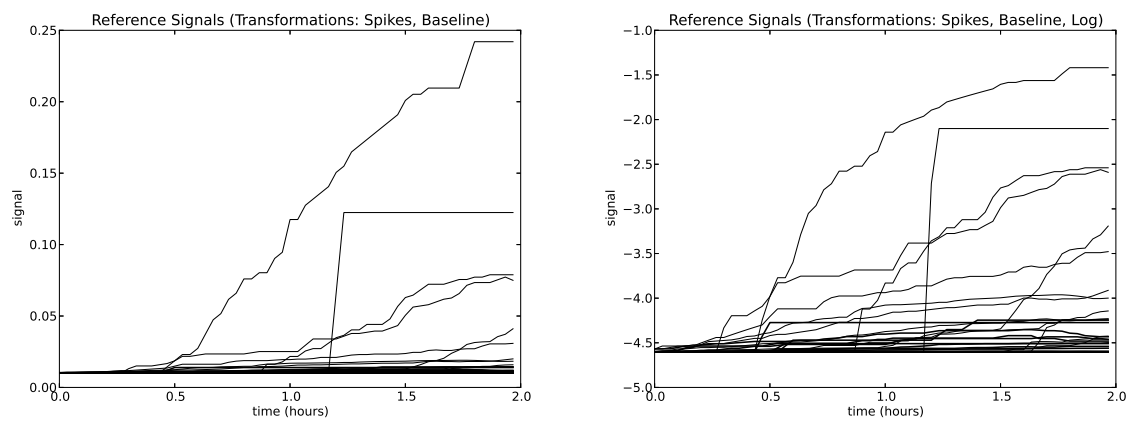


Figure 4-3: Logarithmically scaled reference signals (with spike and baseline normalization previously applied) allow one to make finer-grained distinctions between signals. **Left:** Not logarithmically scaled. **Right:** Logarithmically scaled.

4.3 Experiment

We propose an experiment to measure our algorithm’s performance on two fronts: error rate and relative detection time. We divide the set of topics into a training set and a test set using a 50/50 split. For each topic in the test set, we wish to predict if the topic will trend. If the topic really did trend, we wish to detect it as early as possible relative to the true trend onset while incurring minimal error.

4.3.1 Detection Setup

In principle, to test the detection algorithm, one would step through the signal in the entire sample window for each topic in the test set and report the time of the first detection, or that there were no detections. In practice, we take a shortcut to avoid looking through the entire signal based on the following observations about the activity of topics that trended and topics that did not. First, for topics that trended, there is little, if any activity aside from that surrounding the true onset of the trend. In the rare event that a detection is made very far from the true onset, it is reasonable to assume that this corresponds to a completely different event involving that topic and we can safely ignore it. Thus, the only part of the signal worth looking at is the signal within some time window from the true onset of the trend. Second, topics that did not trend exhibit relatively stationary activity. That is, the signal usually looks roughly the same over the entire sample window. Therefore, it is reasonable to perform detection only on a piece of the signal as an approximation to the true detection performance.

We perform detection over a window of $2N_{obs}$ samples — twice the length of a reference signal. For convenience and future use, we define this in terms of hours.

Definition 6. Let h_{ref} be the number of hours corresponding to N_{ref} samples. At 2 minutes per sample, h_{ref} is given by $N_{ref}/30$.

For test topics that have trended, we do detection on the window spanning $2h_{ref}$ hours centered at the true trend onset. For topics that did not trend, we randomly

choose a window of the desired size. Note that, although this seems to require *a priori* knowledge of whether the test topic ever trended or not, this is only a consequence of the shortcut we take to not do detection over the entire sample window.

4.3.2 Parameter Exploration and Trials

We explore all combinations of the following ranges of parameters, excluding parameter settings that are incompatible (e.g. $N_{obs} > N_{ref}$). For each combination, we conducted 5 random trials.

- γ : 0.1, 1, 10.
- N_{obs} : 10, 80, 115, 150.
- N_{smooth} : 10, 80, 115, 150.
- h_{ref} : 3, 5, 7, 9.
- D_{req} : 1, 3, 5.
- θ : 0.65, 1, 3.